

An aerial photograph of a city square. In the center is a large, ornate, multi-story building with a prominent blue-domed tower and clock faces. The building is surrounded by other city buildings and parking lots. In the foreground, a large crowd of people is gathered on a green lawn, many sitting on blankets. The scene is captured from a high angle, looking down on the square and the surrounding urban area.

# GPU Architectures and Basic CUDA Programming

Tong Shu

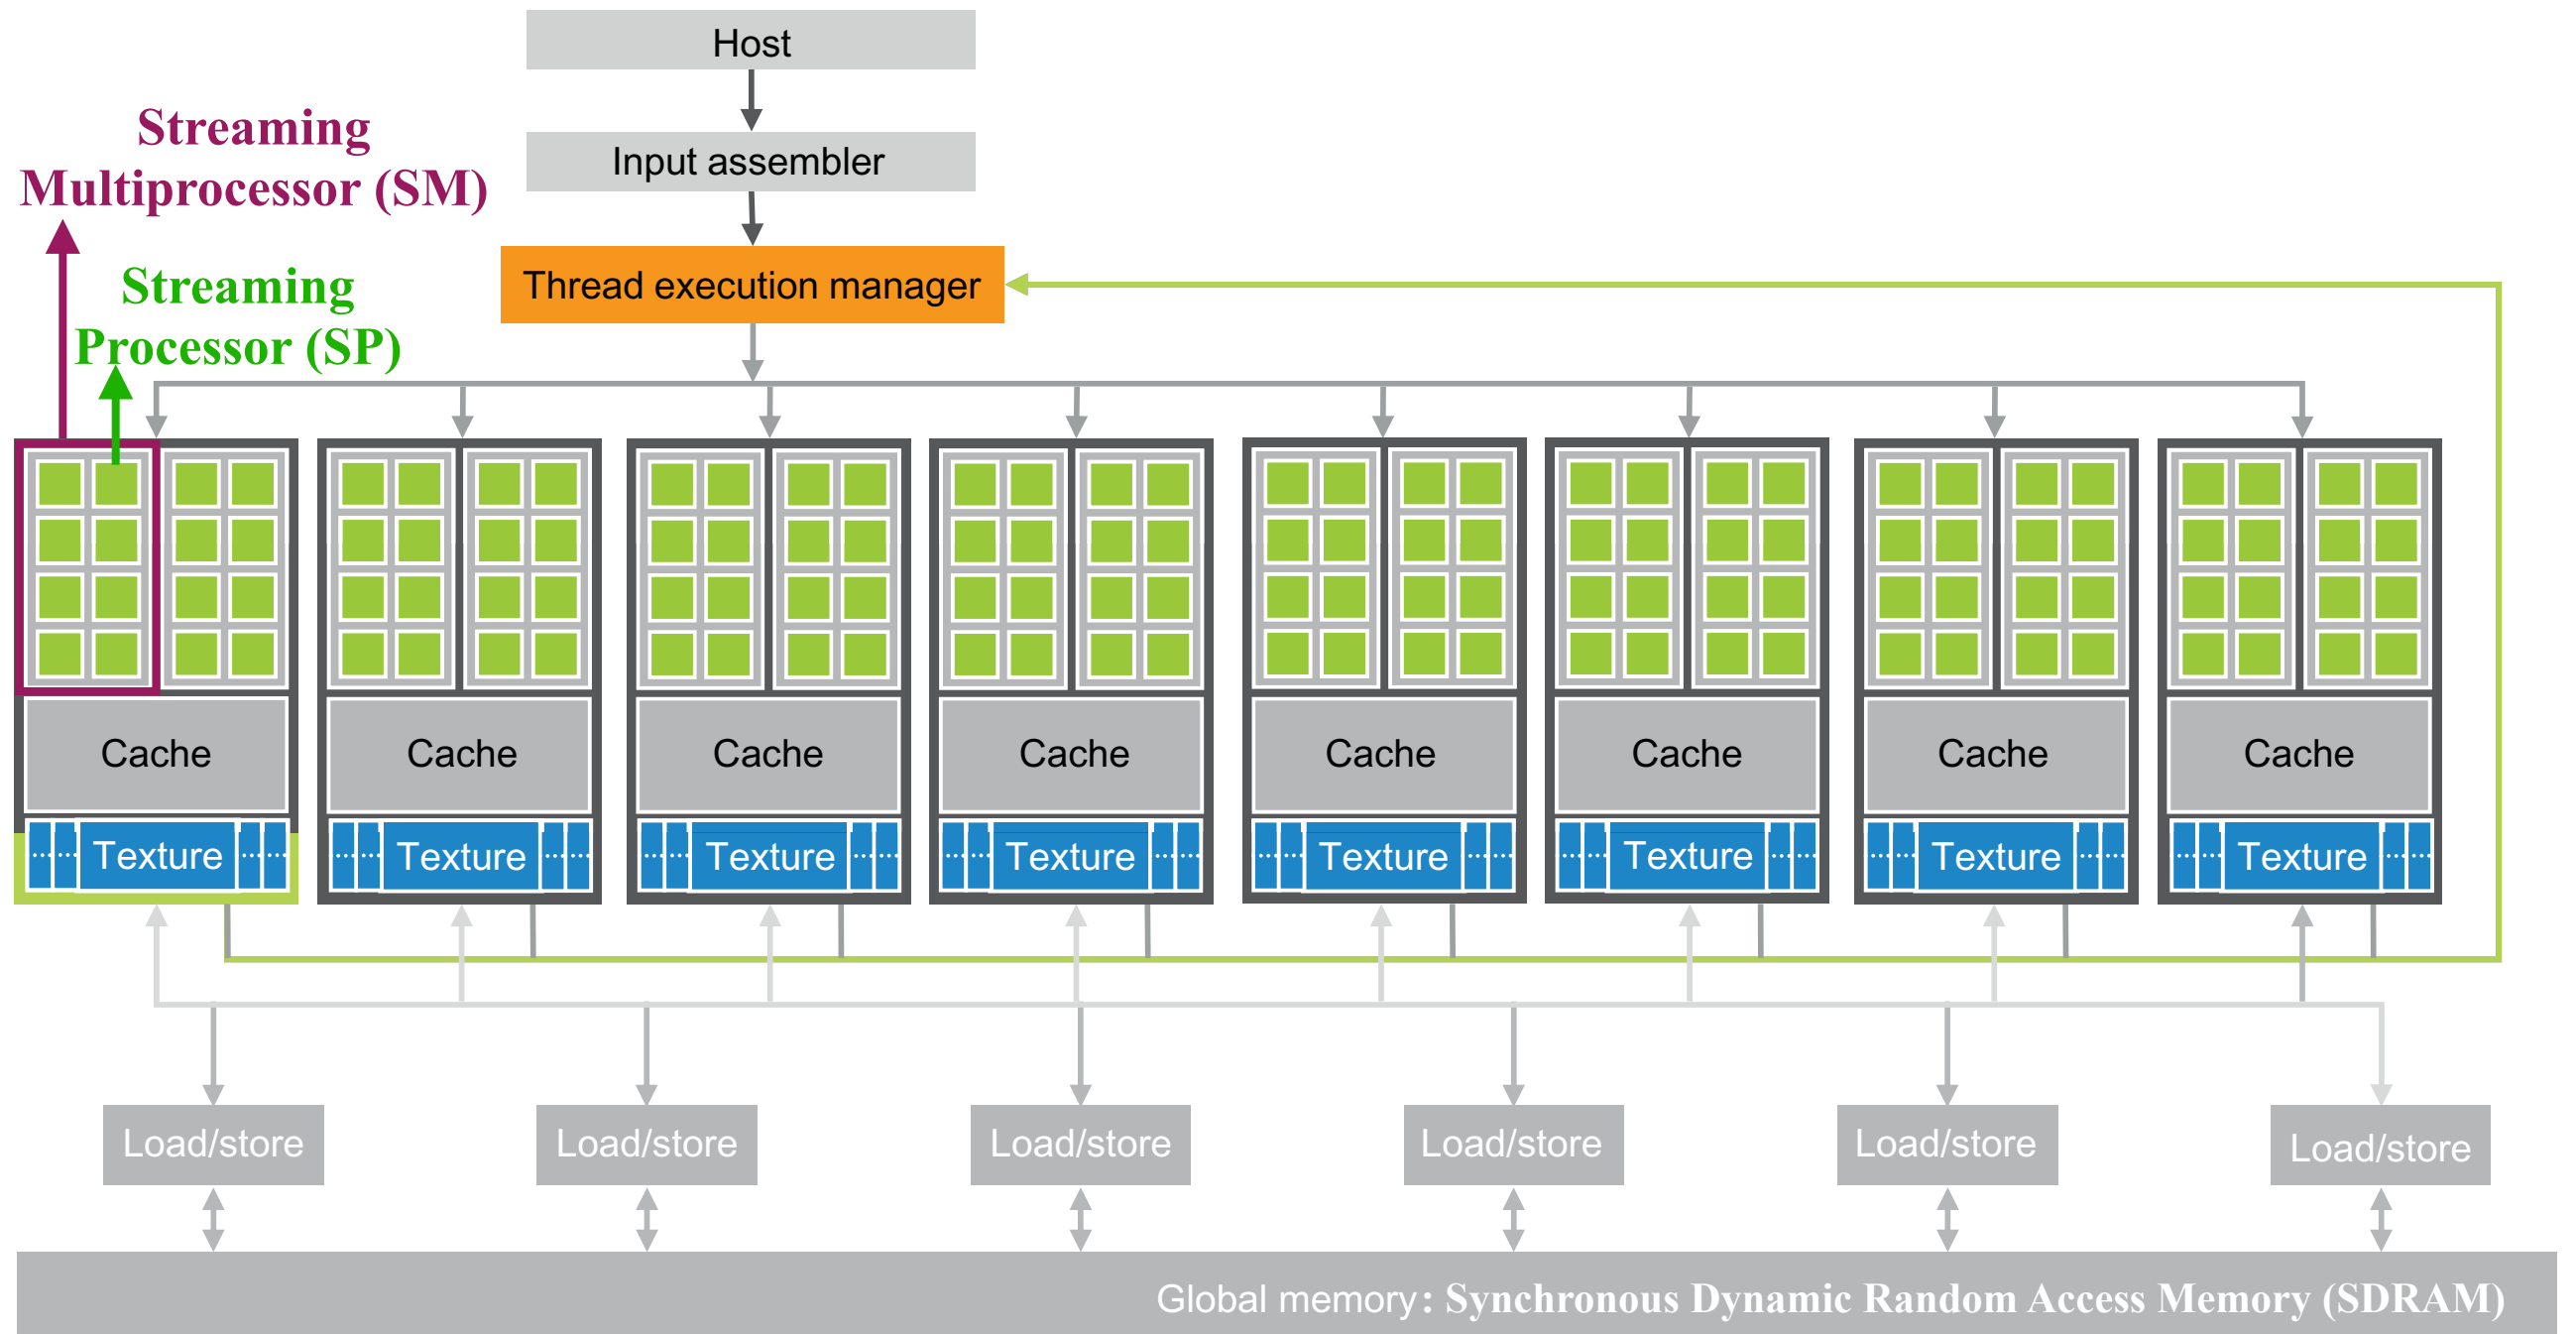
Day 3 — Part 2  
DL-GPU Workshop 2023

07/26/2023



# **Architectures of State-of-the-Art GPUs**

# Architecture of a CUDA-capable GPU

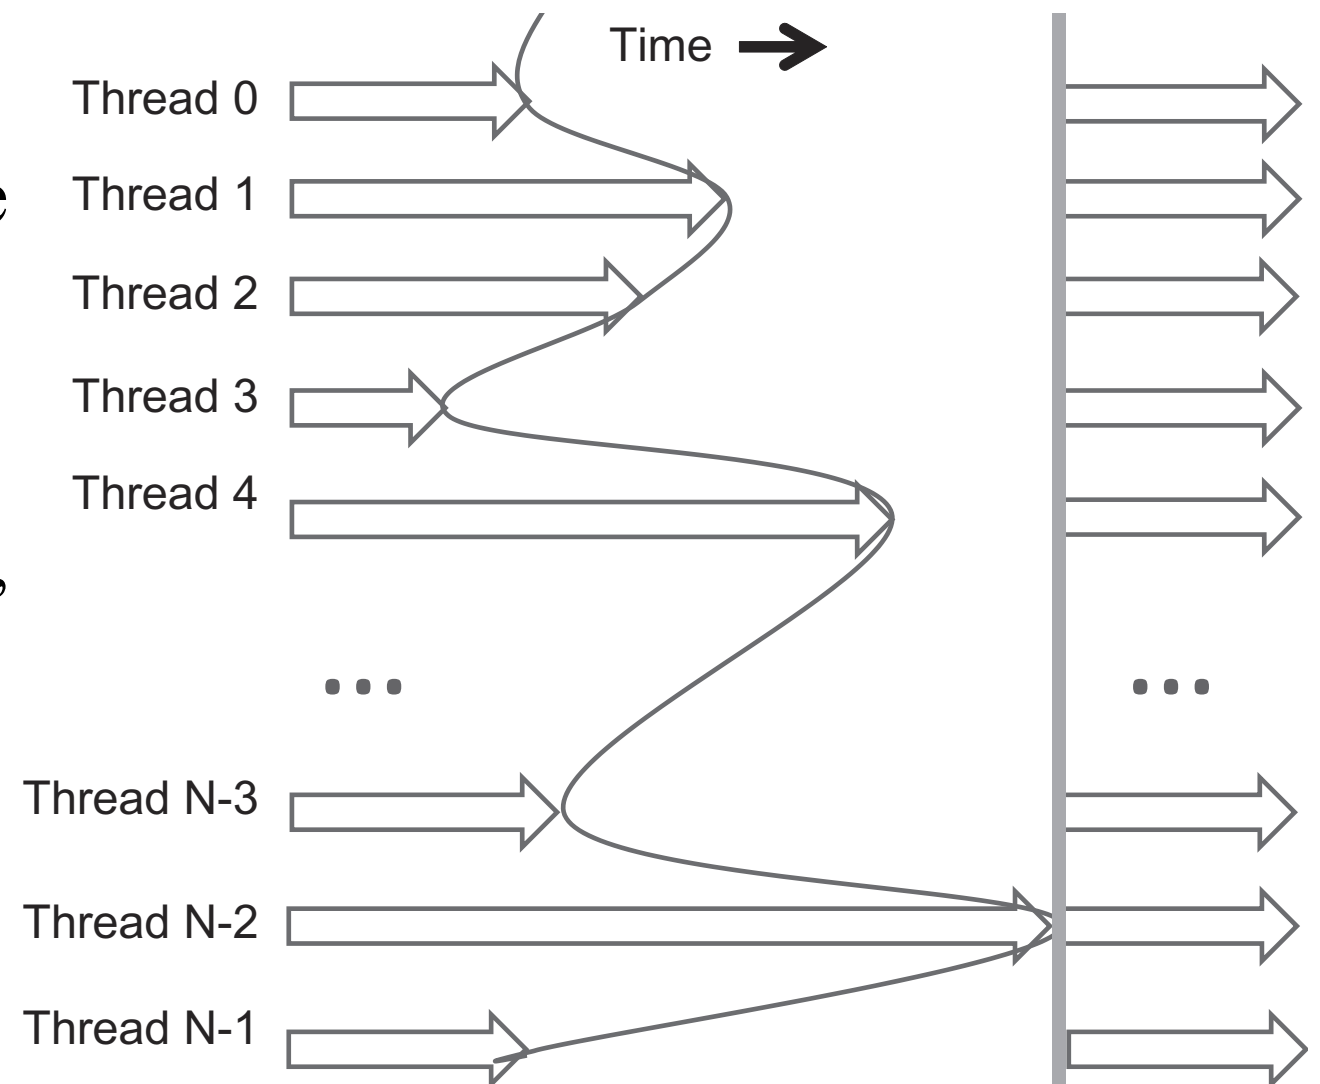


GPU SDRAMs differ from system DRAMs on the CPU motherboard because they are essentially the frame buffer memory used for graphics.

- For graphics applications: hold video images and texture information for 3D rendering.
- For computing: function as very high-bandwidth off-chip memory.

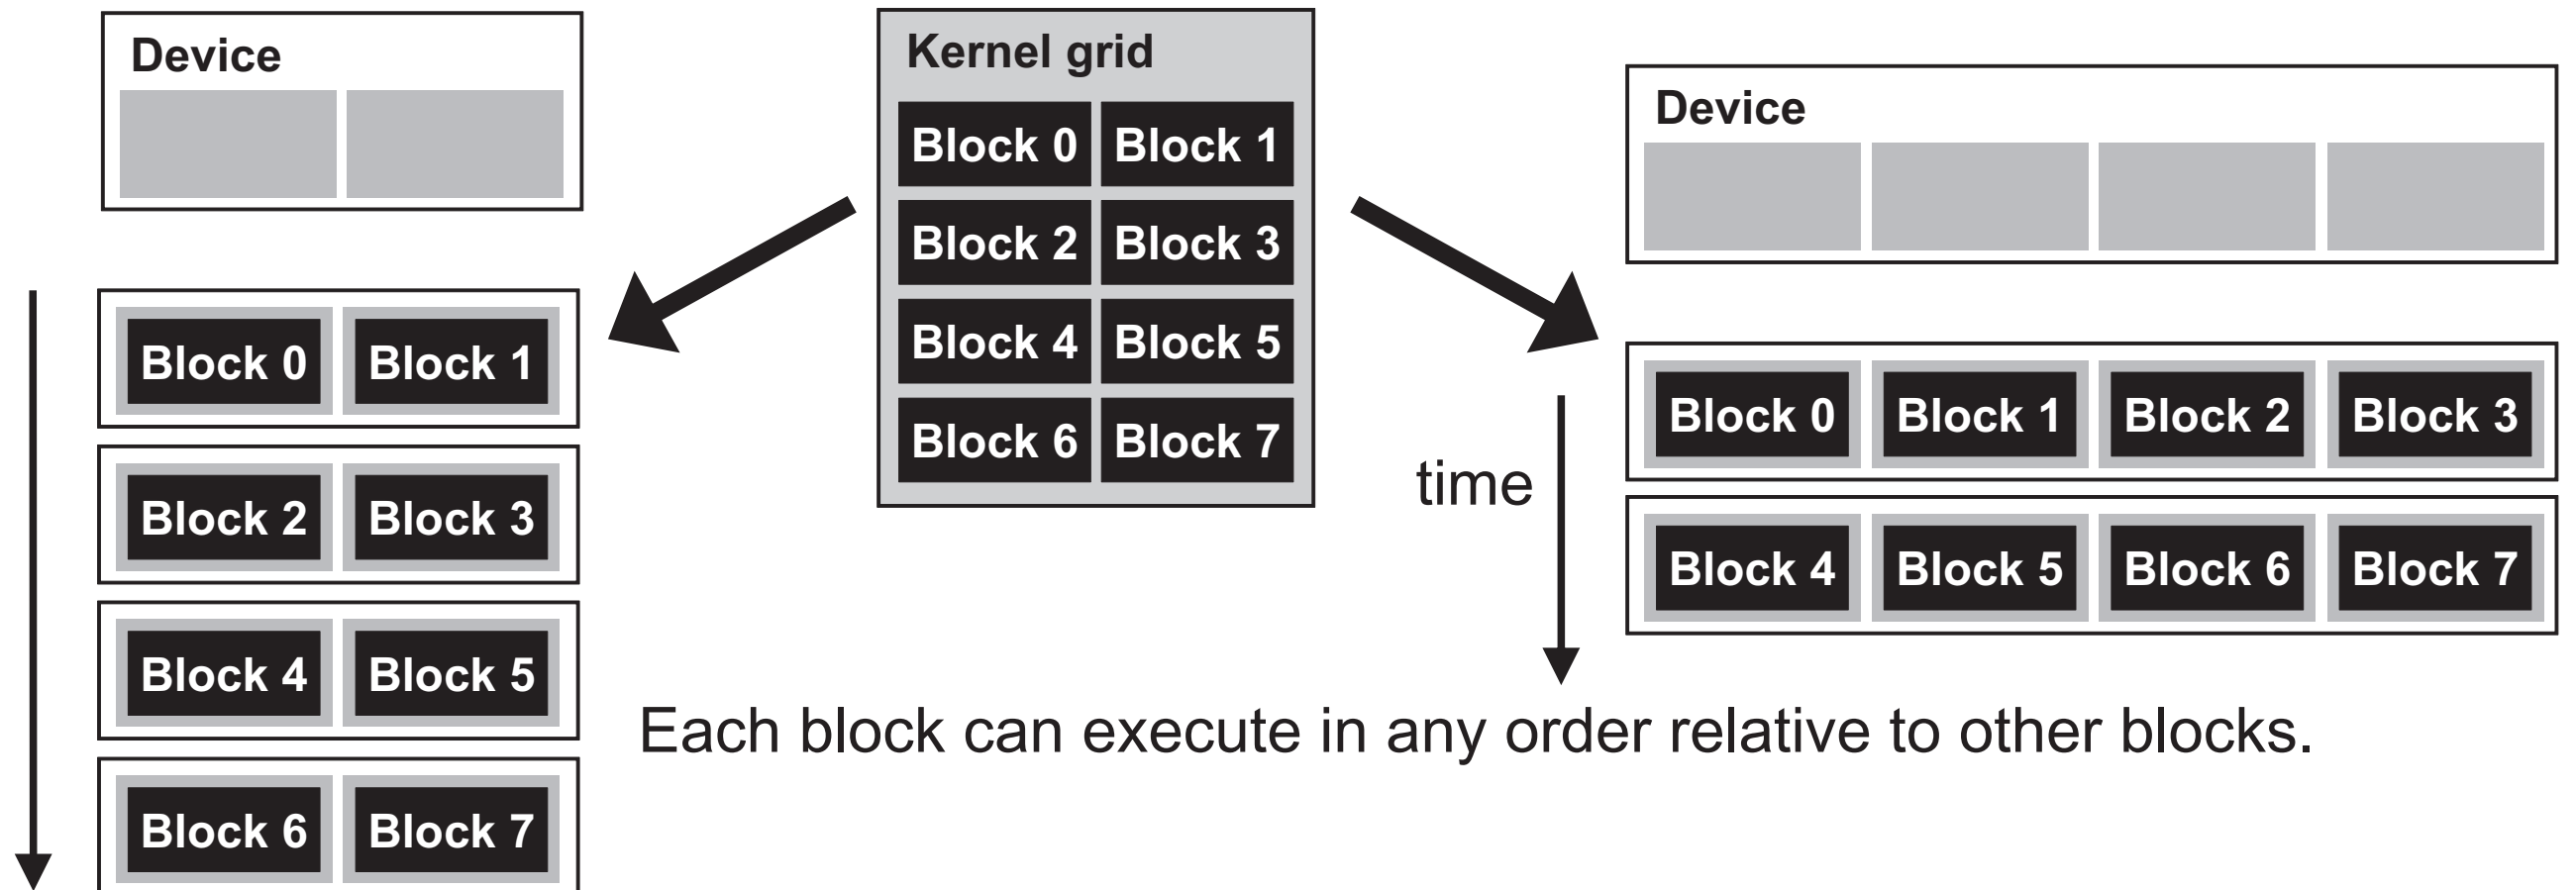
# Thread Synchronization within a Block

- When a thread calls the barrier synchronization function `__syncthreads()`, it will be held at the calling location until every thread in the same block reaches the location.
  - ▶ For an if-then-else statement, if each path has a `__syncthreads()` statement, either all threads in a block execute the then-path or all of them execute the else-path.
- CUDA runtime systems assign the same execution resources to all threads in a block as a unit.
  - ▶ Ensure the temporal proximity of all threads in a block
  - ▶ Prevent excessive or indefinite waiting time during barrier synchronization



# Transparent Scalability


- The ability to execute the same application code on hardware with different numbers of execution resources is referred to as transparent scalability.



- To allow a kernel to maintain transparent scalability, the simple method for threads in different blocks to synchronize with each other is to terminate the kernel and start a new kernel for the activities after the synchronization point.

# Resource Assignment

- Multiple blocks can be assigned to each streaming multiprocessor (SM).



- Limitations on
  - ▶ The number of SMs
  - ▶ The number of blocks that can be assigned to each SM
  - ▶ The number of blocks that can be actively executing in a CUDA device/GPU
  - ▶ The number of threads that can be assigned to an SM
- The runtime system maintains a list of blocks that need to execute and assigns new blocks to SMs when previously assigned blocks complete execution.

# CUDA C API: Querying the Number of Devices

- `__host__ __device__ cudaError_t cudaGetDeviceCount ( int* count )`
  - ▶ Parameters
    - ◎ `count`: Returns the number of compute-capable devices
  - ▶ Returns
    - ◎ `cudaSuccess = 0`
  - ▶ Description
    - ◎ Returns in `*count` the number of devices with compute capability greater or equal to 2.0 that are available for execution.
  - ▶ Example

```
int count;  
cudaGetDeviceCount(&count);
```

# CUDA C API: Querying Device Properties

- `__host__ cudaError_t cudaGetDeviceProperties ( cudaDeviceProp* prop, int device )`
  - ▶ Parameters
    - ◎ `prop`: properties for the specified device
    - ◎ `device`: device number to get properties for
      - ◆ The CUDA runtime numbers all available devices in the system from 0 to count-1, where count is returned by function `cudaGetDeviceCount`.
  - ▶ Returns
    - ◎ `cudaSuccess = 0`
    - ◎ `cudaErrorInvalidDevice = 101`: The device ordinal supplied by the user does not correspond to a valid CUDA device.
  - ▶ Description: returns in `*prop` the properties of the device `device`.
  - ▶ Example

```
cudaDeviceProp prop;
for (int i = 0; i < count; i++) {
    cudaGetDeviceProperties(&prop, i);
    // Decide if device has sufficient resources and capabilities
}
```

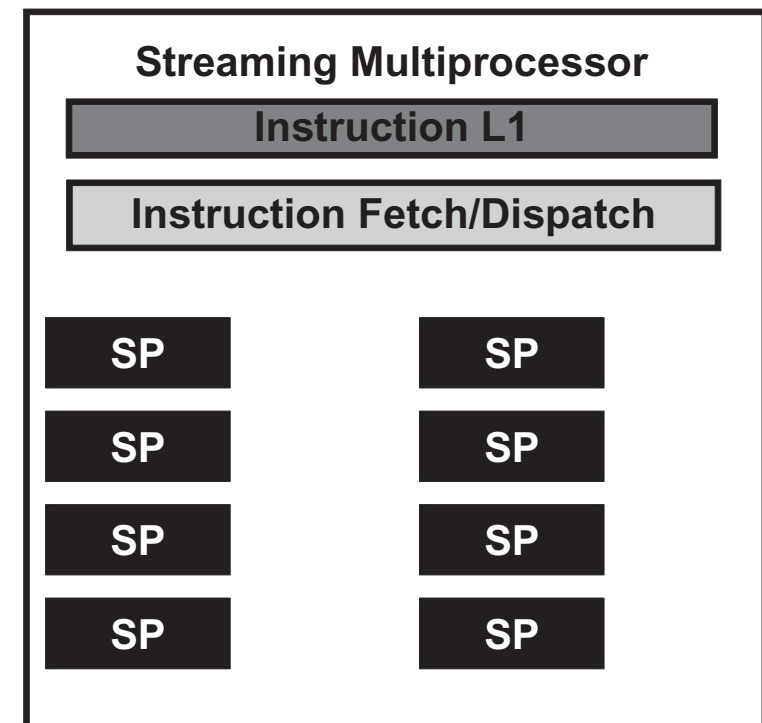



# CUDA C Built-in Type: `cudaDeviceProp`

- C struct type with fields representing the properties of a CUDA device
  - ▶ Field `prop.maxThreadsPerBlock` with the type of `int`: the maximum number of threads allowed in a block in the queried device
  - ▶ Fields `prop.maxThreadsDim[0]`, `prop.maxThreadsDim[1]`, and `prop.maxThreadsDim[2]` with the type of `int`: the maximum number of threads allowed along the x, y, and z dimensions of a block
  - ▶ Fields `prop.maxGridSize[0]`, `prop.maxGridSize[1]`, and `prop.maxGridSize[2]` with the type of `int`: the maximum number of blocks allowed along the x, y, and z dimensions of a grid
  - ▶ Field `prop.warpSize` with the type of `int`: the warp size in threads
  - ▶ Field `prop.multiProcessorCount` with the type of `int`: the number of streaming multiprocessors on the device
  - ▶ Field `prop.clockRate` with the type of `int`: clock frequency in kilohertz

# Warp

- A block is divided into 32 thread units, called warps.
- The size of warps is implementation-specific.
- The warp is the unit of thread scheduling in SMs.
- Each warp consists of 32 threads of consecutive threadIdx values.
- An SM execute all threads in a warp following the Single Instruction, Multiple Data (SIMD) model, i.e., at any instant in time, one instruction is fetched and executed for all threads in the warp.
- The hardware can execute instructions for a small subset of all warps in the SM.




# Thread Scheduling and Latency Tolerance

- Latency tolerance/hiding: the mechanism of filling the latency time of operations with work from other threads
  - ▶ When an instruction to be executed by a warp needs to wait for the result of a previously initiated long-latency operation, another resident warp that is no longer waiting for results is selected for execution.
- Zero-overhead thread scheduling
  - ▶ The selection of ready warps for execution avoids introducing idle or wasted time into the execution timeline



# NVIDIA H100 GPU on new SXM5 Module




- <https://developer.nvidia.com/blog/nvidia-hopper-architecture-in-depth>

# Hopper Architecture Features of Nvidia H100 GPU

- The fourth-generation tensor cores perform faster matrix computations than ever before on an even broader array of AI and HPC tasks.
- A new transformer engine enables H100 to deliver up to 9x faster AI training and up to 30x faster AI inference speedups on large language models compared to Nvidia A100.
- The new NVLink network interconnect enables GPU-to-GPU communication among up to 256 GPUs across multiple compute nodes.
- Secure MIG partitions the GPU into isolated, right-size instances to maximize quality of service (QoS) for smaller workloads.

# H100 v.s. A100 in HPC and AI







# H100 SM Feature

- H100 FP16 Tensor Core has 3x throughput compared to A100 FP16 Tensor Core

A100 FP16




H100 FP16



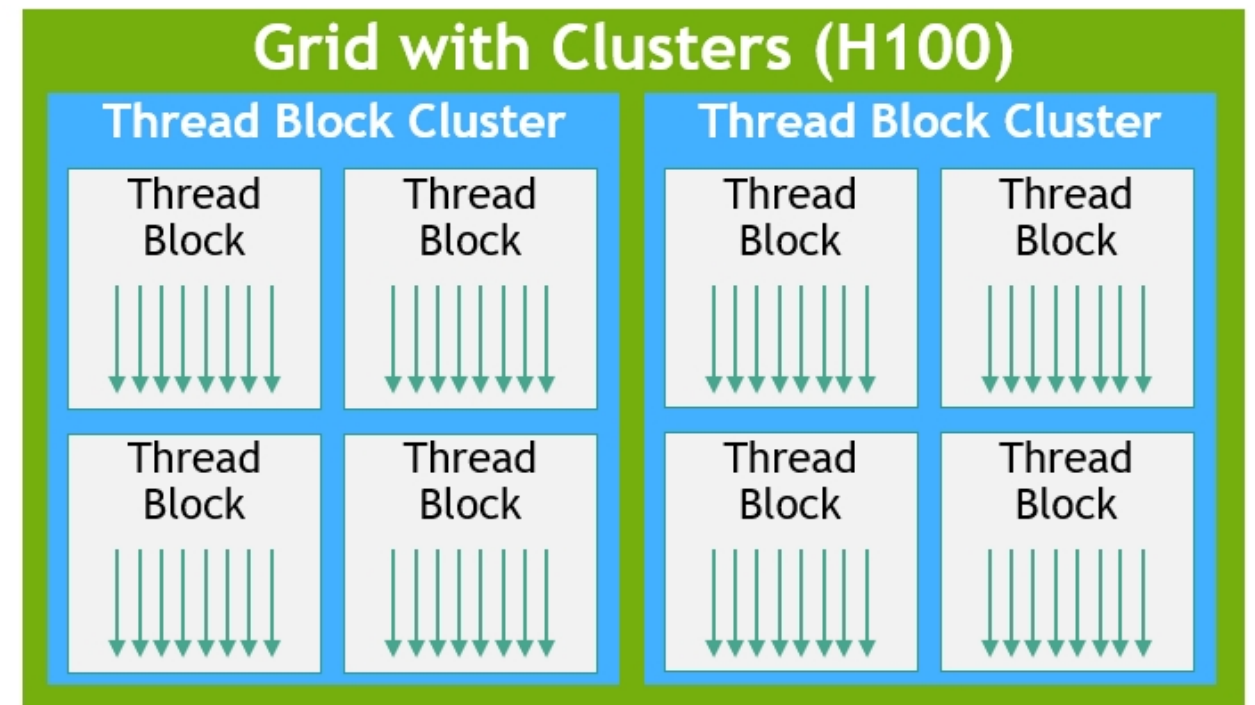
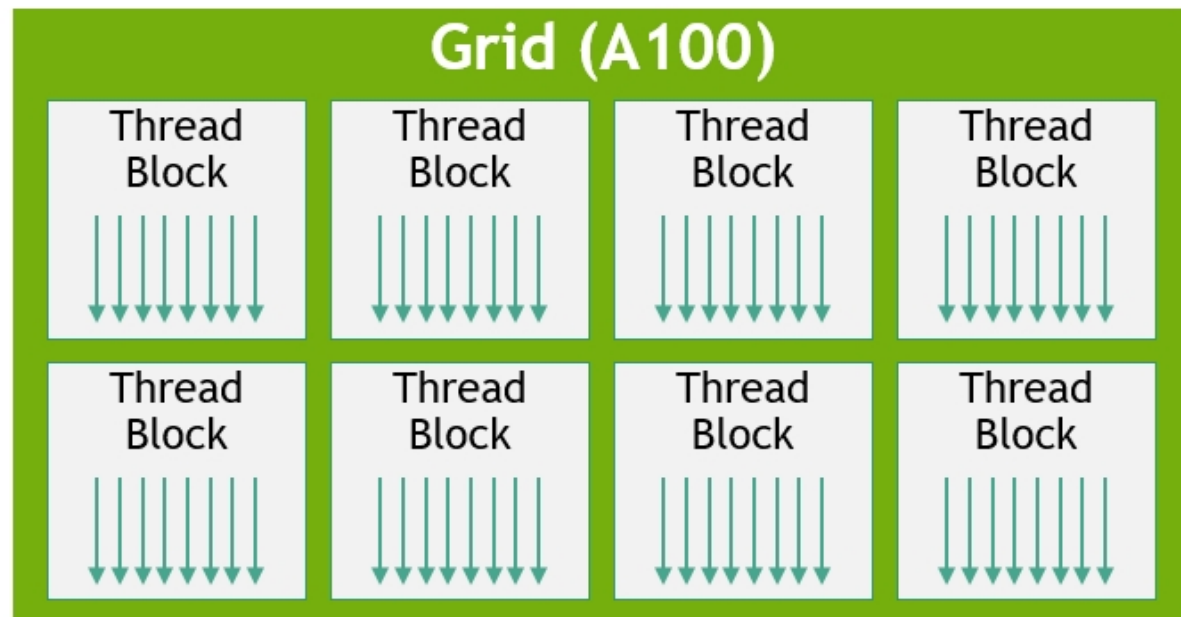
# H100 SM Feature

- H100 compute improvement
  - ▶ 132 SMs provide a 22% SM count increase over the A100 108 SMs
  - ▶ Each of the H100 SMs is 2x faster, due to the fourth-generation Tensor Core.
  - ▶ Within each Tensor Core, the new FP8 format and associated transformer engine provide another 2x improvement.
  - ▶ Increased clock frequencies in H100 deliver another approximately 1.3x performance improvement.



# H100 SM Feature

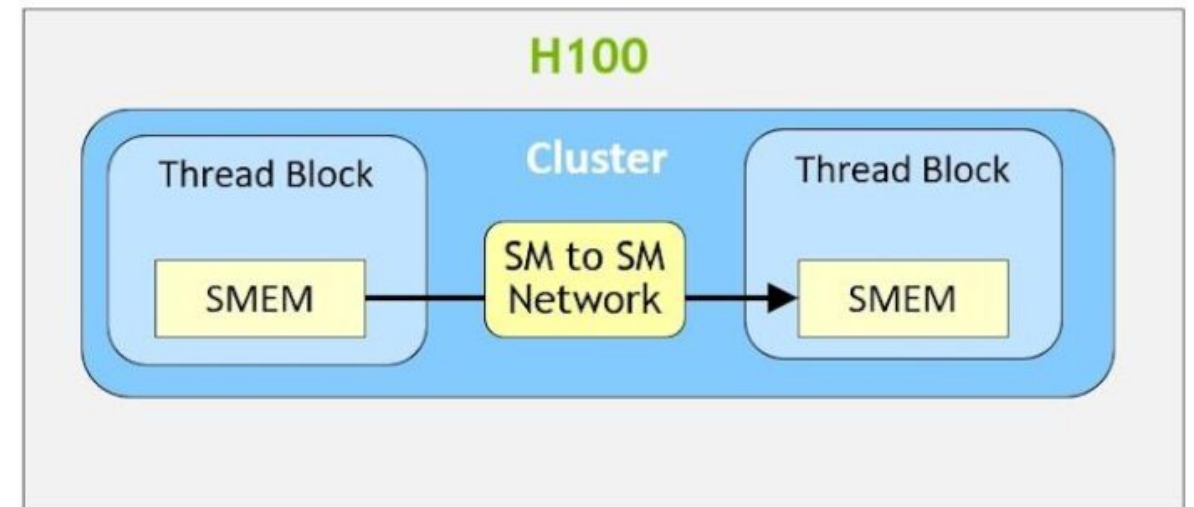
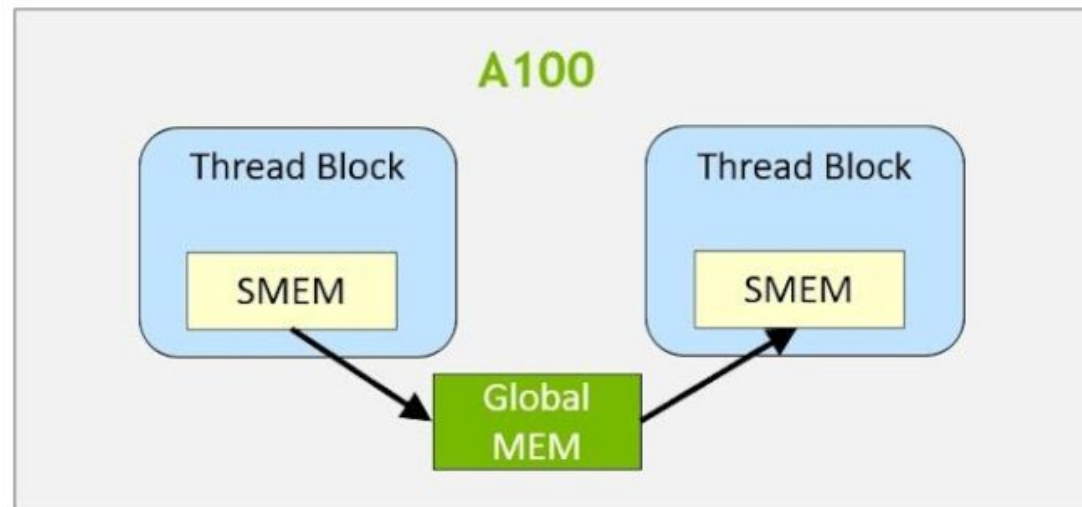
- Thread block clusters and grids with clusters
  - ▶ A cluster is a group of thread blocks that are guaranteed to be concurrently scheduled onto a group of SMs.
  - ▶ Clusters have hardware-accelerated barriers and new memory access collaboration capabilities.
  - ▶ A dedicated SM-to-SM network for SMs in a GPU processing cluster provides fast data sharing between threads in a cluster.
  - ▶ Cluster capabilities can be leveraged from the CUDA cooperative\_groups API.





# H100 SM Feature

- Distributed shared memory
  - ▶ Thread-block-to-thread-block data exchange (A100 vs. H100 with clusters)



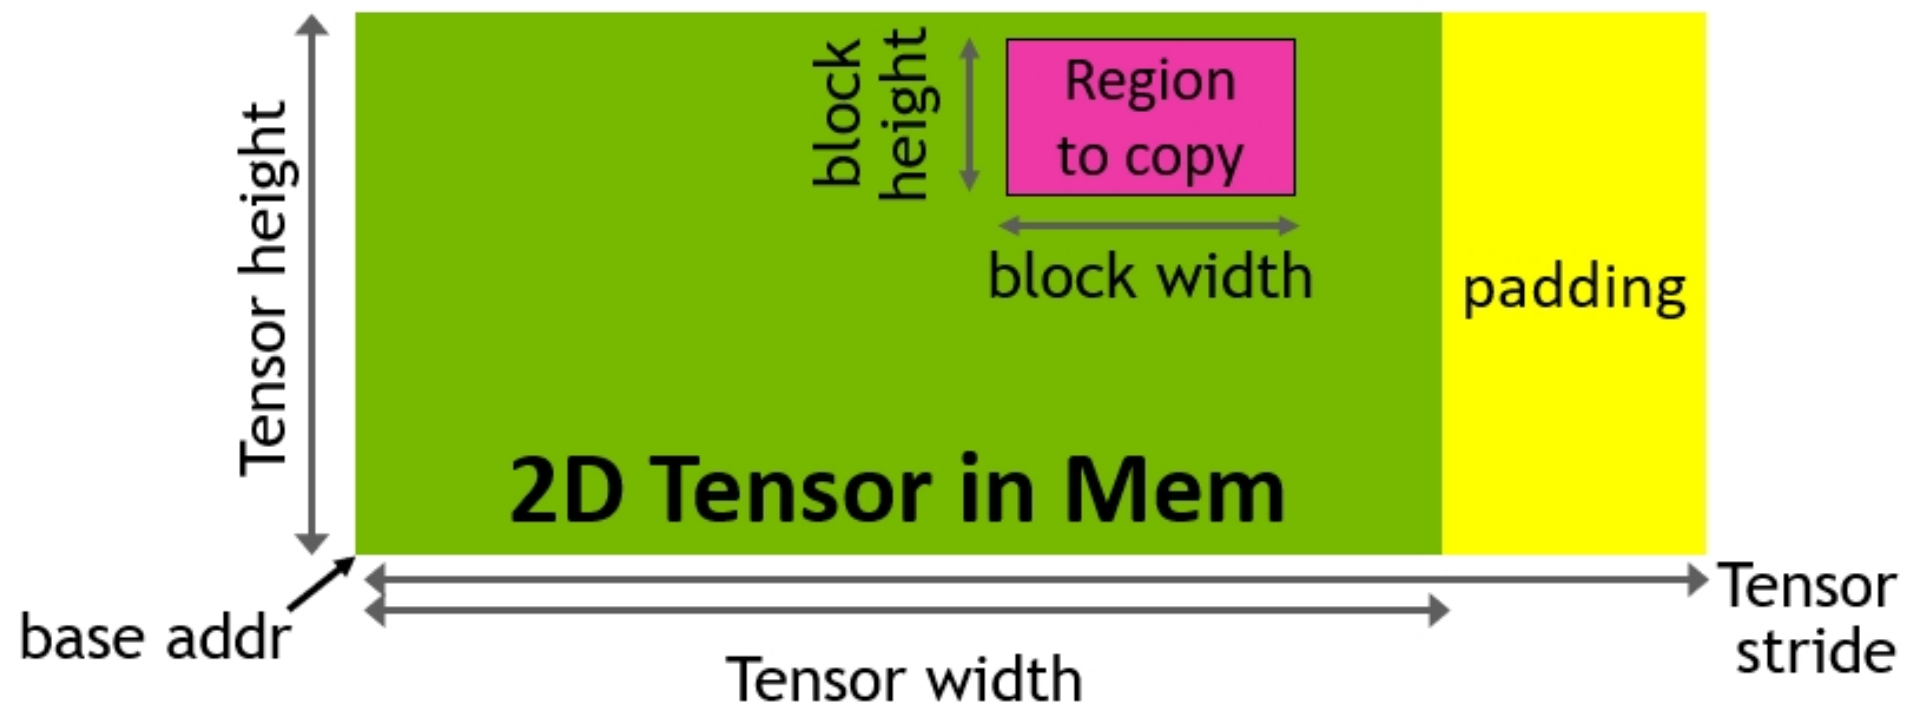
- ▶ Cluster vs. non-cluster performance comparisons

## Cluster Performance



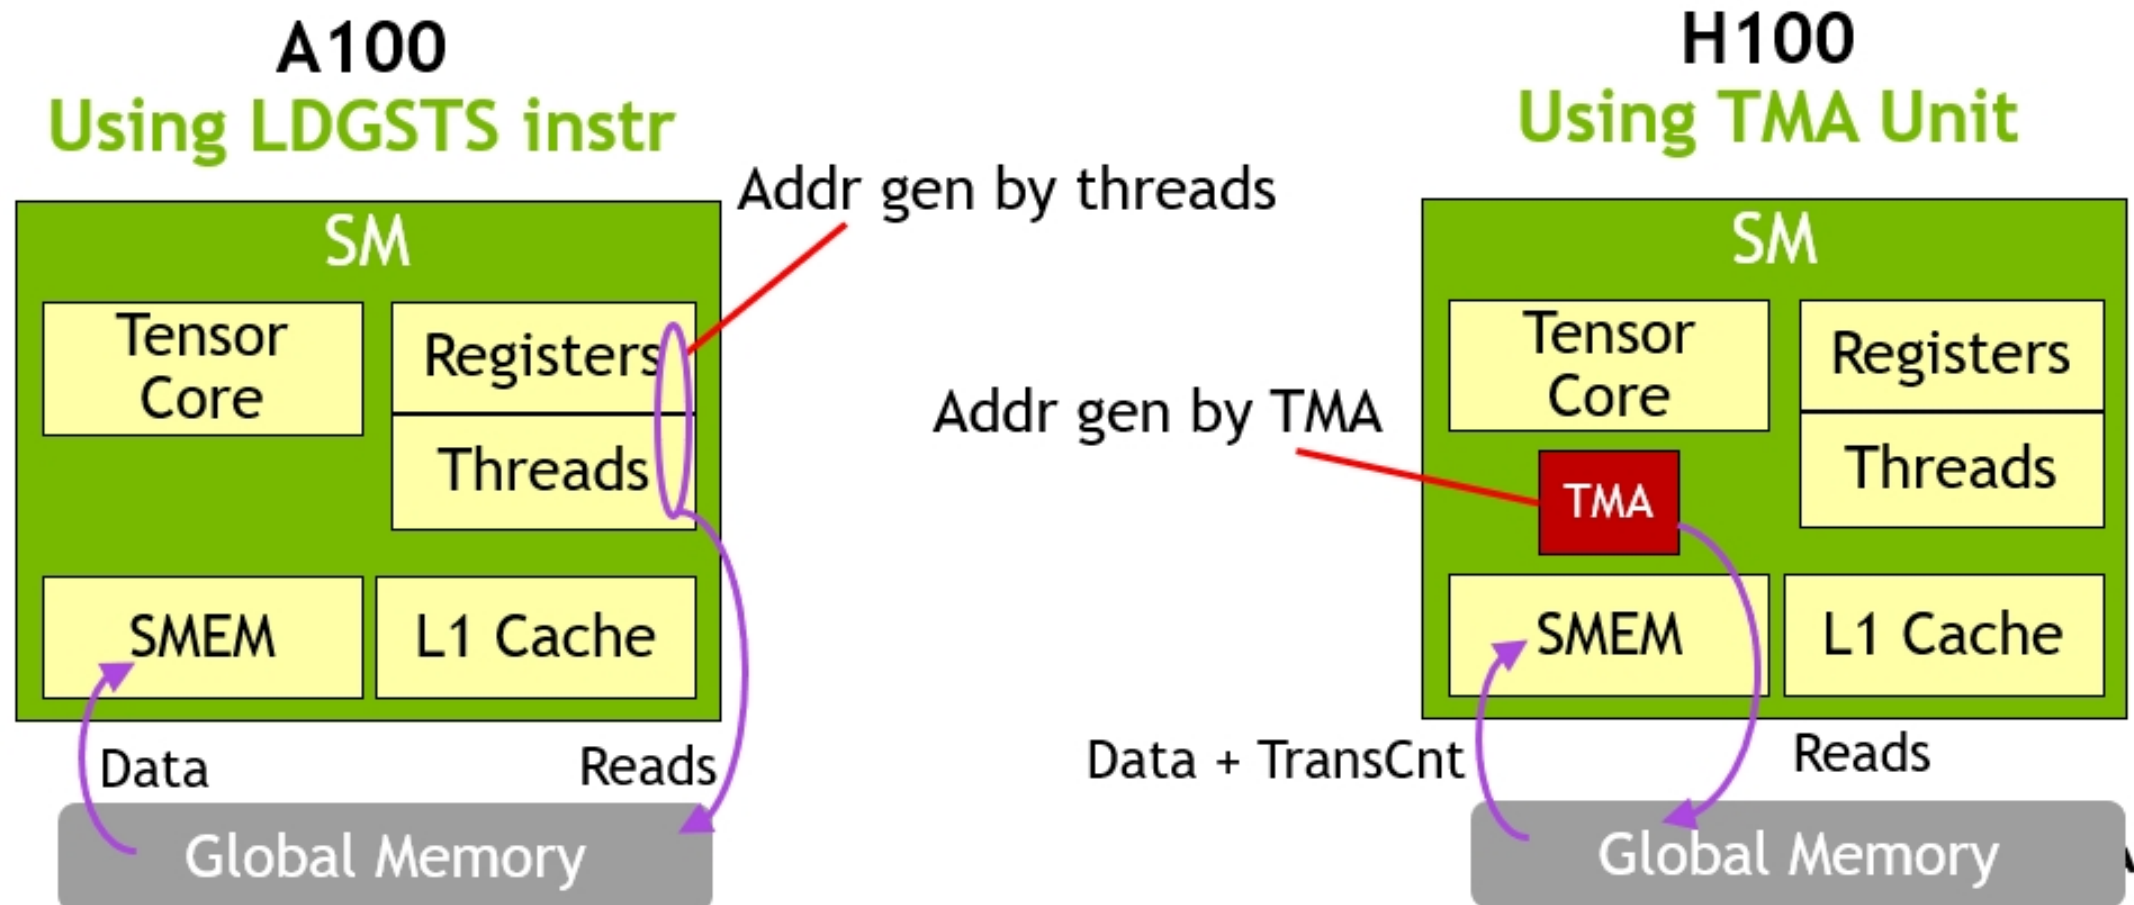
# H100 SM Feature

- Tensor memory accelerator: efficiently transfers large blocks of data between global memory and shared memory.
  - ▶ TMA address generation through a copy descriptor that specifies data transfers using tensor dimensions and block coordinates instead of per-element addressing



# Tensor Memory Accelerator


- A single thread in a warp is elected to issue an asynchronous TMA operation (`cuda::memcpy_async`) to copy a tensor.
- Multiple threads can wait on a shared memory-based asynchronous barrier `cuda::barrier` for the completion of the data transfer.
- Advantage: it frees the threads to execute other independent work.
- Asynchronous memory copy with TMA on H100 vs. LoadGlobalStoreShared (LDGSTS) on A100




# H100 SM Feature

- Asynchronous barrier in A100
  - ▶ First, threads signal Arrive when they are done producing their portion of the shared data. This Arrive is non-blocking so that the threads are free to execute other independent work.
  - ▶ Eventually, the threads need the data produced by all the other threads. At this point, they do a Wait, which blocks them until every thread has signaled Arrive.
- Asynchronous transaction barrier in H100

## Asynchronous Barrier (from A100)



## Async Transaction Barrier (New on H100)





# H100 SM Feature

- Fourth-generation NVLink
  - ▶ A100 includes 12 third-generation NVLink links to provide 600 GB/sec total bandwidth.
  - ▶ H100 includes 18 fourth-generation NVLink links to provide 900 GB/sec total bandwidth.
- Fourth-generation NVLink Network
  - ▶ H100 introduces the new NVLink Network interconnect, a scalable version of NVLink that enables GPU-to-GPU communication among up to 256 GPUs across multiple compute nodes.
- Third-generation NVSwitch
  - ▶ NVLink Switch System supports up to 256 GPUs. The connected nodes can deliver 57.6 TBs of all-to-all bandwidth and can supply an incredible one exaFLOP of FP8 sparse AI compute.

# **Memory and Data Locality in CUDA Programming**

# Memory and Data Locality

- Compute-to-global-memory-access ratio
  - ▶ The number of floating-point calculation performed for each access to the global memory within a region of a program
- Example
  - ▶ In a high-end device, the global memory bandwidth is around 1 TB/s.
  - ▶  $1000/4 = 250$  giga single-precision operands per second can be loaded.
  - ▶ 250 giga floating-point operations per second (GFLOPS)
  - ▶ 2% of 12 TFLOPS for these high-end devices

```
for (int blurRow = -BLUR_SIZE; blurRow < BLUR_SIZE + 1; ++blurRow) {
    for (int blurCol = -BLUR_SIZE; blurCol < BLUR_SIZE + 1; ++blurCol) {
        int curRow = row + blurRow;
        int curCol = col + blurCol;
        if (curRow >= 0 && curRow <= h - 1 && curCol >= 0 && curCol <= w - 1) {
            sumPixRed += d_in[curRow * w + curCol].r;
            sumPixGreen += d_in[curRow * w + curCol].g;
            sumPixBlue += d_in[curRow * w + curCol].b;
            numPix++;
        }
    }
}
```

Compute-to-global-memory ratio = 1.0

# Compute-bound vs. Memory-bound

- Compute-bound
  - ▶ The time for it to complete a task is determined principally by the speed of the processor.
  - ▶ E.g. Processor utilization is high, perhaps at 100% usage.
- I/O bound
  - ▶ The time it takes to complete a computation is determined principally by the period spent waiting for input/output operations to be completed.
  - ▶ Memory-bound programs
    - ◎ Programs whose execution speed is limited by memory access throughput




# Kernel Function for Matrix Multiplication

// Each thread produces one output matrix element.


```
__global__ void matMulKernel(float* d_C, float* d_A, float* d_B, int n) {  
    // Calculate the column index of the element in d_C  
    int col = blockIdx.x * blockDim.x + threadIdx.x;  
    // Calculate the row index of the element in d_C  
    int row = blockIdx.y * blockDim.y + threadIdx.y;  
  
    if (col < n && row < n) {  
        float dotProc = 0;  
        for (int k = 0; k < n; k++) {  
            dotProc += d_A[row * n + k] * d_B[k * n + col];  
        }  
        d_C[row * n + col] = dotProc;  
    }  
}
```

Compute-to-global-  
memory ratio = 1.0



# CUDA Device Memory Model

- Device code can
  - ▶ Read and write per-thread on-chip registers
    - ◎ Bandwidth of registers is at least two orders of magnitude higher than that of the global memory
  - ▶ Read and write per-block on-chip shared memory
  - ▶ Read and write per-thread off-chip local memory (cached)
  - ▶ Read only per-grid off-chip constant memory (cached)
  - ▶ Read and write per-grid off-chip global memory (DRAM)



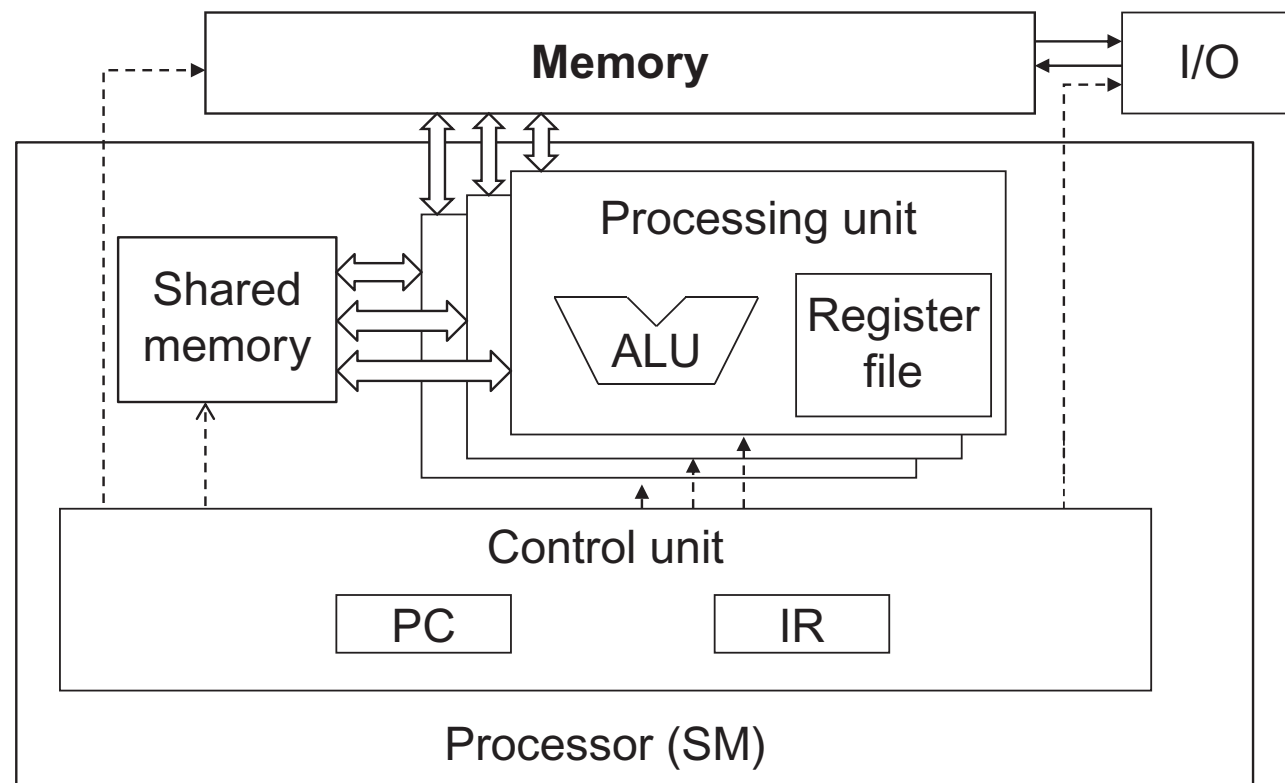
- Host code can
  - ▶ Transfer data to/from per-grid global and constant memories

# Fewer Instructions for Access to Registers than Global Memory

- Arithmetic instructions have “built-in” register operands.
- Example: `fadd r1, r2, r3`
  - ▶ r2 and r3 are the register numbers that specify the location in the register file where the input operand values can be found.
  - ▶ r1 specifies the location for storing the floating-point addition result value.
- If an operand value is in the global memory, the processor needs to perform a memory load operation to make the operand value available to the arithmetic and logic unit (ALU) first.
- Example: assembly language

```
load r2, r4, offset  
fadd r1, r2, r3
```

- ▶ Add an offset value to the content of r4 to form an address for the operand value.
- ▶ Access the global memory and place the operand value into register r2.



# On-chip Shared Memory is Scratchpad Memory

- Latency and bandwidth
  - ▶ Accessed with much lower latency and much higher throughput than the global memory.
  - ▶ Accessed with longer latency and lower bandwidth than on-chip registers
    - ◎ Because the processor needs to perform a memory load operation when accessing data in the shared memory
- Data sharing
  - ▶ Variables in the shared memory are accessible by all threads in a block
  - ▶ Register data are private to a thread.



# Thread

- Control unit
  - ▶ Maintains a program counter (PC), which contains the memory address of the next instruction to be executed.
  - ▶ Uses the PC to fetch an instruction into the instruction register (IR) in each “instruction cycle”
- Context-switching
  - ▶ Multiple threads can time-share a processor by taking turns to make progress.
  - ▶ The execution of a thread can be suspended and then correctly resumed later by saving and restoring the PC value and the contents of registers/memory.
- Single-Instruction, Multiple-Data design
  - ▶ Multiple processing units share a PC and IR, and allow multiple threads to make simultaneous progress.
  - ▶ All threads make simultaneous progress by executing the same instruction.

# Scope vs. Lifetime

- Scope of a variable (a textual or spatial concept in the compiling period)
  - ▶ Serial programming: the range of statements in which the variable is visible
    - ◎ A variable is visible in a statement if it can be referenced or assigned in that statement.
  - ▶ Parallel programming: the range of threads that can access the variable
    - ◎ A single thread only
    - ◎ All threads of a block
    - ◎ All threads of a grid
- Lifetime of a variable (a temporal concept in run-time)
  - ▶ The portion of the program execution duration when the variable is available for use (has valid memory)
    - ◎ Within a kernel execution
    - ◎ Throughout the entire application

# CUDA Variable Types

Variable declaration	Memory	Scope	Lifetime
Automatic scalar variables	Register	Thread	Kernel
Automatic array variables	Local	Thread	Kernel
<code>__device__ __shared__ int SharedVar;</code>	Shared	Block	Kernel
<code>__device__ __constant__ int ConstVar;</code>	Constant	Grid	Application
<code>__device__ int GlobalVar;</code>	Global	Grid	Application

# CUDA Variable Types

- Automatic scalar variables
  - ▶ Scalar variables: variables that are not arrays or matrices
- Automatic array variables
  - ▶ Rarely used in kernel/device functions
- Shared variables
  - ▶ Reside within a kernel/device function
  - ▶ Hold global memory data heavily used in a kernel execution phase
- Constant variables
  - ▶ Declaration is outside any function body.
  - ▶ Provide input values to kernel functions
  - ▶ Stored in the global memory but cached for efficient access
  - ▶ The total size in an application is limited to 65,536 bytes.
- Global variables
  - ▶ The way to synchronize between threads from different blocks or to ensure data consistency across threads when accessing global memory is by terminating the current kernel execution.
  - ▶ Pass information from one kernel invocation to another kernel invocation



# *Question & Answer*