

---

# DL-GPU PROJECT 2:

Exploration of TPU Architectures for  
Optimized Transformer Performance in  
Image Detection of Drainage Crossings

Presented by Aikaterini Panteleaki, Amir Nazeri, and Denys Godwin

---

---

# PROJECT DESCRIPTION

- As high-resolution digital elevation models (HRDEMs) continue to improve the mapping of hydrographic features, the challenge of accurately identifying under-road drainage structures like culverts remains significant.
  - Utilizing the Transformer architecture, renowned for its ability to handle complex visual data through self-attention mechanisms, this project refines image analysis techniques for environmental science.
  - By optimizing the performance of Transformer models through targeted TPU architecture exploration using the Scalesim tool , this project minimizes latency, compute time, and cost while maintaining high accuracy when using the model for inference.
-

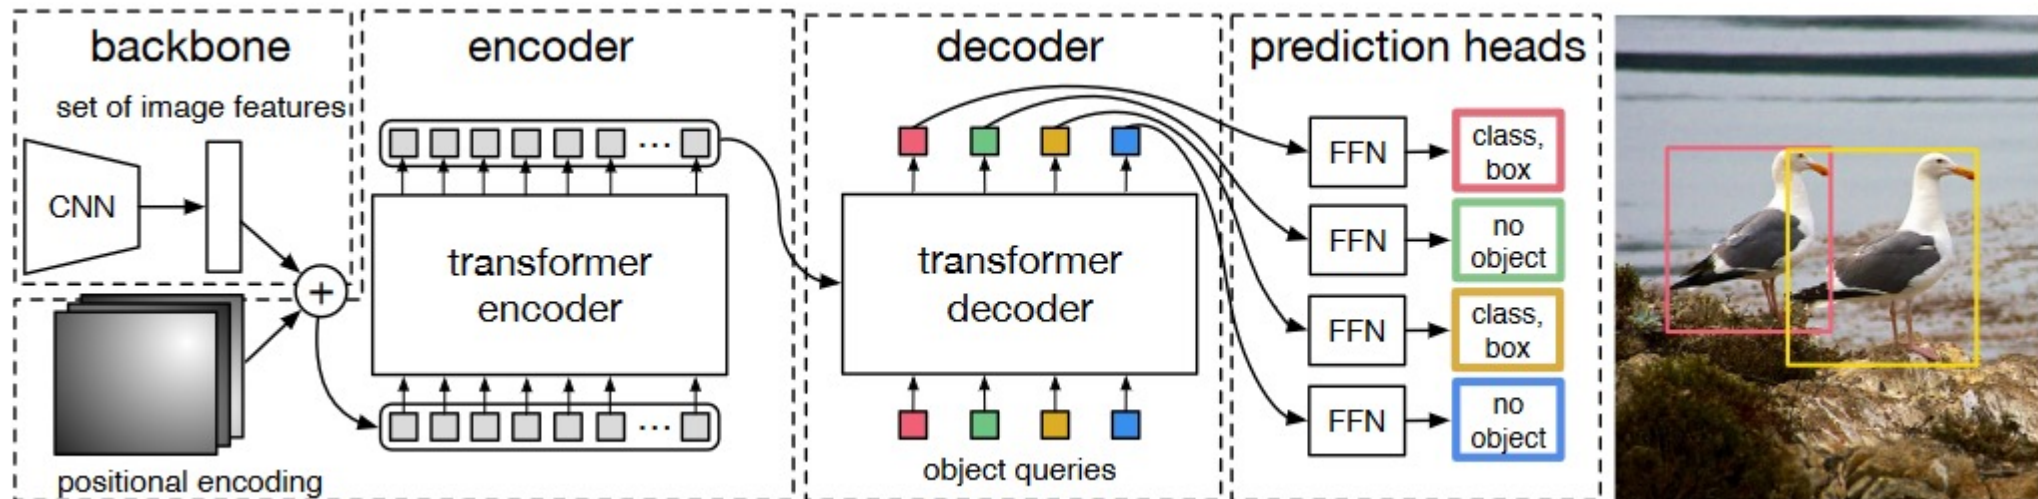
---

# THE MODEL

---

# DETECTION -TRANSFORMER (DETR)

- Introduced by Facebook AI in 2020.



---

# DETR FEATURES

- Key advantages of DETR over traditional object detection:
- End-to-end training pipeline: No need for hand-crafted component in Faster RCNN or YOLO
- Transformer architecture: transformer component leverages self-attention mechanism to capture complex and long-dependencies.
- Simplified architecture: No RPN or non-maximum suppression (NMS). Abstracted by Bipartite matching loss.
- Easier Scalability: Embedded transformer accelerates processing of large data.

---

# COCO DATASET

- COCO2017 dataset, containing 118k training images and 5k validation images.
- 7 instances per image on average and up to 63 instances in a single training image.
- Includes annotated with 80 object categories and 5 captions describing the scene..

```
"categories": [  
  {  
    "id": 1,  
    "name": "car",  
    "supercategory": "vehicle"  
  },  
  {  
    "id": 2,  
    "name": "truck",  
    "supercategory": "vehicle"  
  }  
]
```

```
"images": [  
  {  
    "id": 1,  
    "width": 640,  
    "height": 480,  
    "file_name": "000000397133.jpg",  
    "license": 1,  
    "flickr_url": "https://www.flickr.com/photos/adrianrosebrock/397133",  
    "coco_url": "http://images.cocodataset.org/val2017/000000397133.jpg",  
    "date_captured": "2013-11-14 17:02:52"  
  },  
]
```

---

# DETR RESULTS ON COCO

- Inference on validation set:

```
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.420
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.624
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.442
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.205
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.458
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.611
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.333
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.533
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.574
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.312
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.628
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.805
```

---

# MODELS' COMPARISON

Model	GFLOPS/FPS	#params	AP	AP <sub>50</sub>	AP <sub>75</sub>	AP <sub>S</sub>	AP <sub>M</sub>	AP <sub>L</sub>
Faster RCNN-DC5	320/16	166M	39.0	60.5	42.3	21.4	43.5	52.5
Faster RCNN-FPN	180/26	42M	40.2	61.0	43.8	24.2	43.5	52.0
Faster RCNN-R101-FPN	246/20	60M	42.0	62.5	45.9	25.2	45.6	54.6
Faster RCNN-DC5+	320/16	166M	41.1	61.4	44.3	22.9	45.9	55.0
Faster RCNN-FPN+	180/26	42M	42.0	62.1	45.5	26.6	45.4	53.4
Faster RCNN-R101-FPN+	246/20	60M	44.0	63.9	<b>47.8</b>	<b>27.2</b>	48.1	56.0
DETR	86/28	41M	42.0	62.4	44.2	20.5	45.8	61.1
DETR-DC5	187/12	41M	43.3	63.1	45.9	22.5	47.3	61.1
DETR-R101	152/20	60M	43.5	63.8	46.4	21.9	48.0	61.8
DETR-DC5-R101	253/10	60M	<b>44.9</b>	<b>64.7</b>	47.7	23.7	<b>49.5</b>	<b>62.3</b>



---

# DATASET PREPARATION

---

---

# THE DATASET

- The original dataset as shared is formatted to match Pascal VOC standard for object detection
- Each single-channel 800x800 .tif file has a corresponding .xml file containing the annotations for the image
- .tif files represent LiDAR-derived High-Resolution Digital Elevation Model (HRDEM) data
  - Elevation represented as units above sea level
  - Dataset stats:

Mean	Standard Deviation	Max	Min	Mean Elevation Difference
3,960.79	10,892.09	47,075	3.48	114.2667

# PRE-PROCESSING THE DATASET

- Images with invalid pixel values (either  $-1E38$  or anomalously high values) are filtered in `prepare_dataset.ipynb`:

```
df.at[index, 'valid'] = False if (np.max(data) - np.min(data)) > 10000 else True
```

- Images are converted to relative elevation from absolute elevation with code from `get_normalization.ipynb`:

```
min_value = data.min()
new_data = data - min_value
```

- Then, the mean and standard deviation is extracted from the entire dataset with the code on the right to generate new stats:

Mean	Standard Deviation	Max	Min	Mean Elevation Difference
55.82	185.58	3581	0	114.36

```
psum = torch.tensor([0.0])
psum_sq = torch.tensor([0.0])
min = torch.tensor

index = 0
# Loop through images
for inputs in tqdm(image_loader):
    psum += inputs.sum()
    psum_sq += (inputs ** 2).sum()
    index += 1

# pixel count
count = len(tif_files) * 800 * 800

# mean and STD
total_mean = psum / count
total_var = (psum_sq / count) - (total_mean ** 2)
total_std = torch.sqrt(total_var)
```

---


# MODIFYING THE DATASET TO FIT THE MODEL

- DETR requires a COCO-formatted dataset
- Image chips were split randomly into train (70%), validate (20%), and test (10%) directories
- Image ids and properties were stored in a COCO-formatted json
- Annotation bounding boxes were extracted from the xml files and compiled into the COCO-formatted JSON, where each annotation is associated with an image id
  - This is implemented in `prepare_dataset.ipynb`

---

# MODIFYING THE MODEL TO FIT THE DATASET

- The default COCO dataloader imported from pycocotools uses PIL to load imagery in RGB format
  - This leads to degradation of resolution (see example, right)
  - Solved by writing a custom dataloader which uses rasterio to load .tif files, repeat along the first axis for 3 channels, and convert to torch.Tensor format.



---

# CHALLENGES IN ADAPTING DETR TO HRDEM DATA

- The data augmentations and transforms expect input in a PIL Image.Image format, and therefore will need to be modified to expect torch.Tensor input
- COCO formatting requires precision in conversion
  - Issues arose around using str for int fields and vice versa
  - Issues arose around image ids not matching correctly with corresponding annotations

---

# MODEL PERFORMANCE

---

---


# TEST DESCRIPTION

- Two preliminary model tests were run, each using a batch size of 8 and training for 100 Epochs
    - Each training took 13 hours and 20 minutes with a maximum GPU utilization of 7.6GB on one Quadro RTX 5000
  - Model 1 is the stock 90-class model pretrained on the COCO dataset
    - "Drainage Culvert" is assigned to category 1, replacing "Person"
  - Model 2 is a custom binary classifier model with randomly initialized weights
-



# MODEL 1: PRETRAINED PERFORMANCE





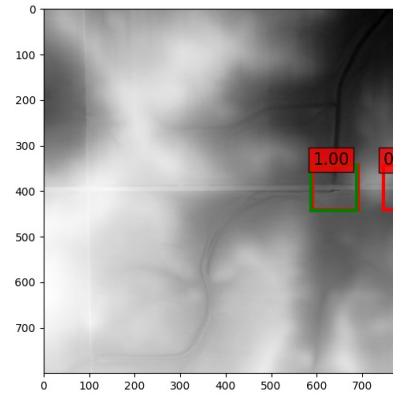
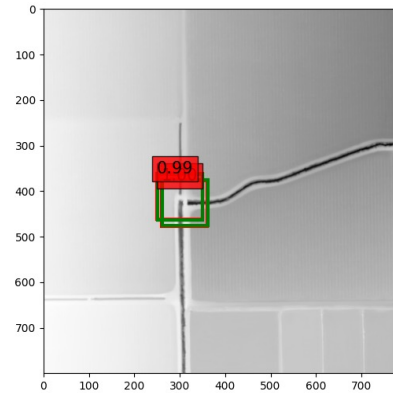
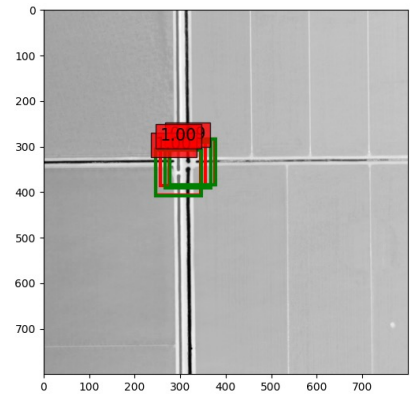

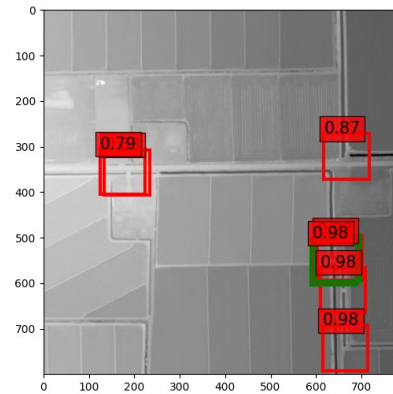
---



---

## MODEL 1 VISUALIZATIONS: PRETRAINED

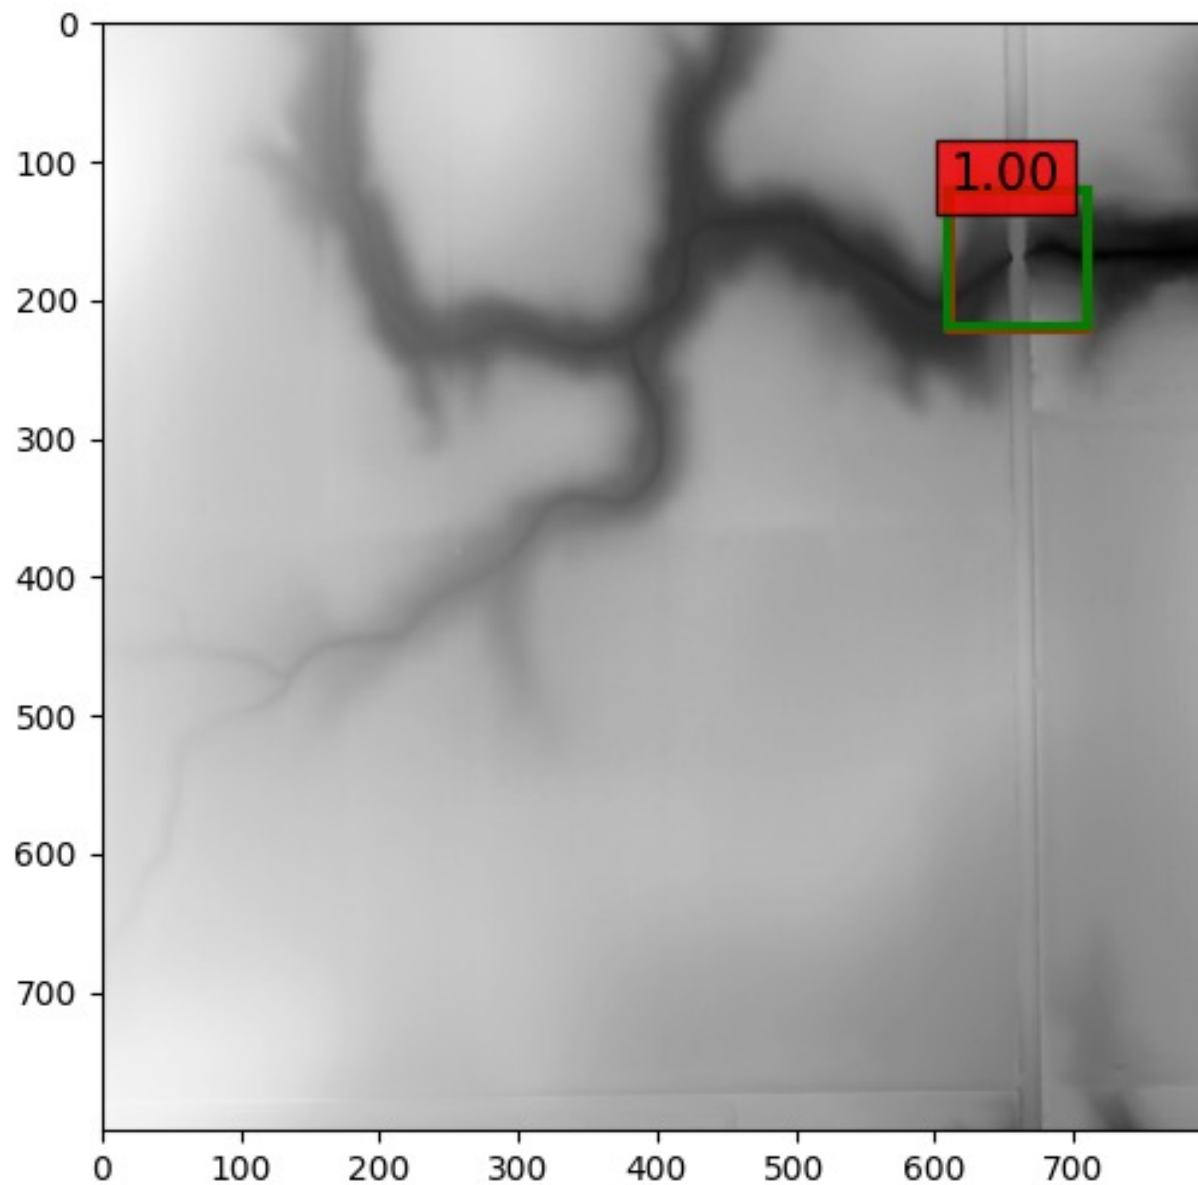
- All test set visualizations for Model 1 may be viewed at <https://rb.gy/1gczzw>
  - To avoid cherry-picking, this is every 70th visualization in the test set (idx 0, 70, 140 etc.)
  - Only predictions with over 70% confidence are visualized
  - Of these visualizations:
    - 6 match the target perfectly
    - None contain a miss
    - 1 matches the target AND finds an unambiguous culvert that the human classifier missed
    - 2 match the target and find ambiguous candidates for a culvert
- 



---

# MODEL 1 VISUALIZATIONS: PRETRAINED

Case 1: Single Easy Object




---

## MODEL 1 VISUALIZATIONS: PRETRAINED

### Case 2: Multiple Objects

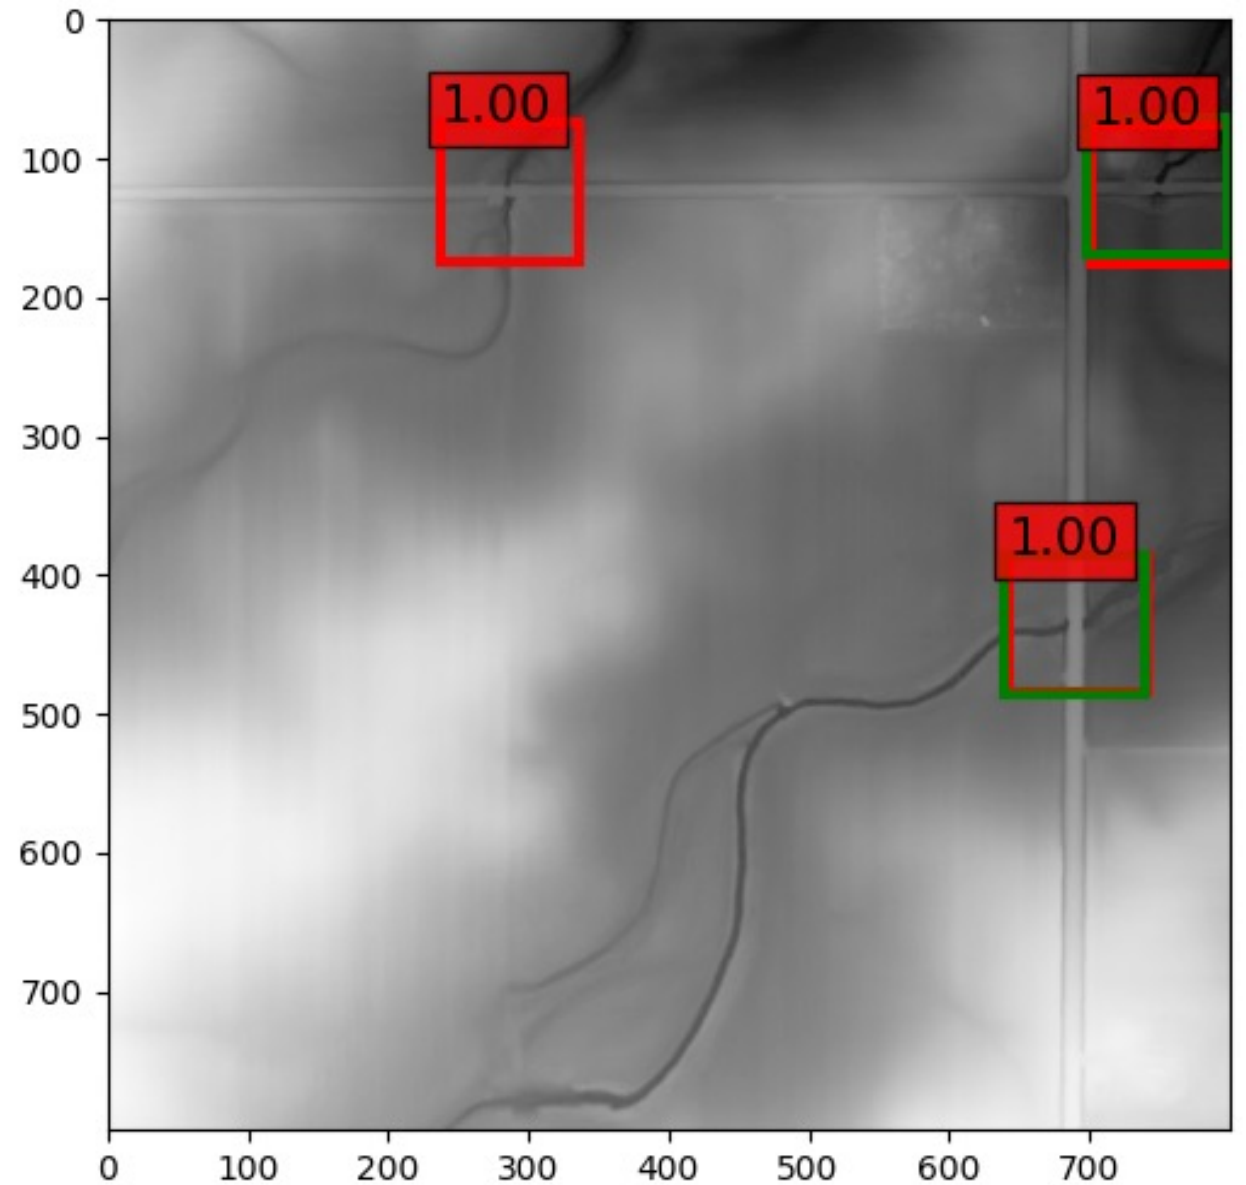
- Multiple drainage culverts are detected for each target box – this is likely because this is a complex system of multiple culverts



---

## MODEL 1 VISUALIZATIONS: PRETRAINED

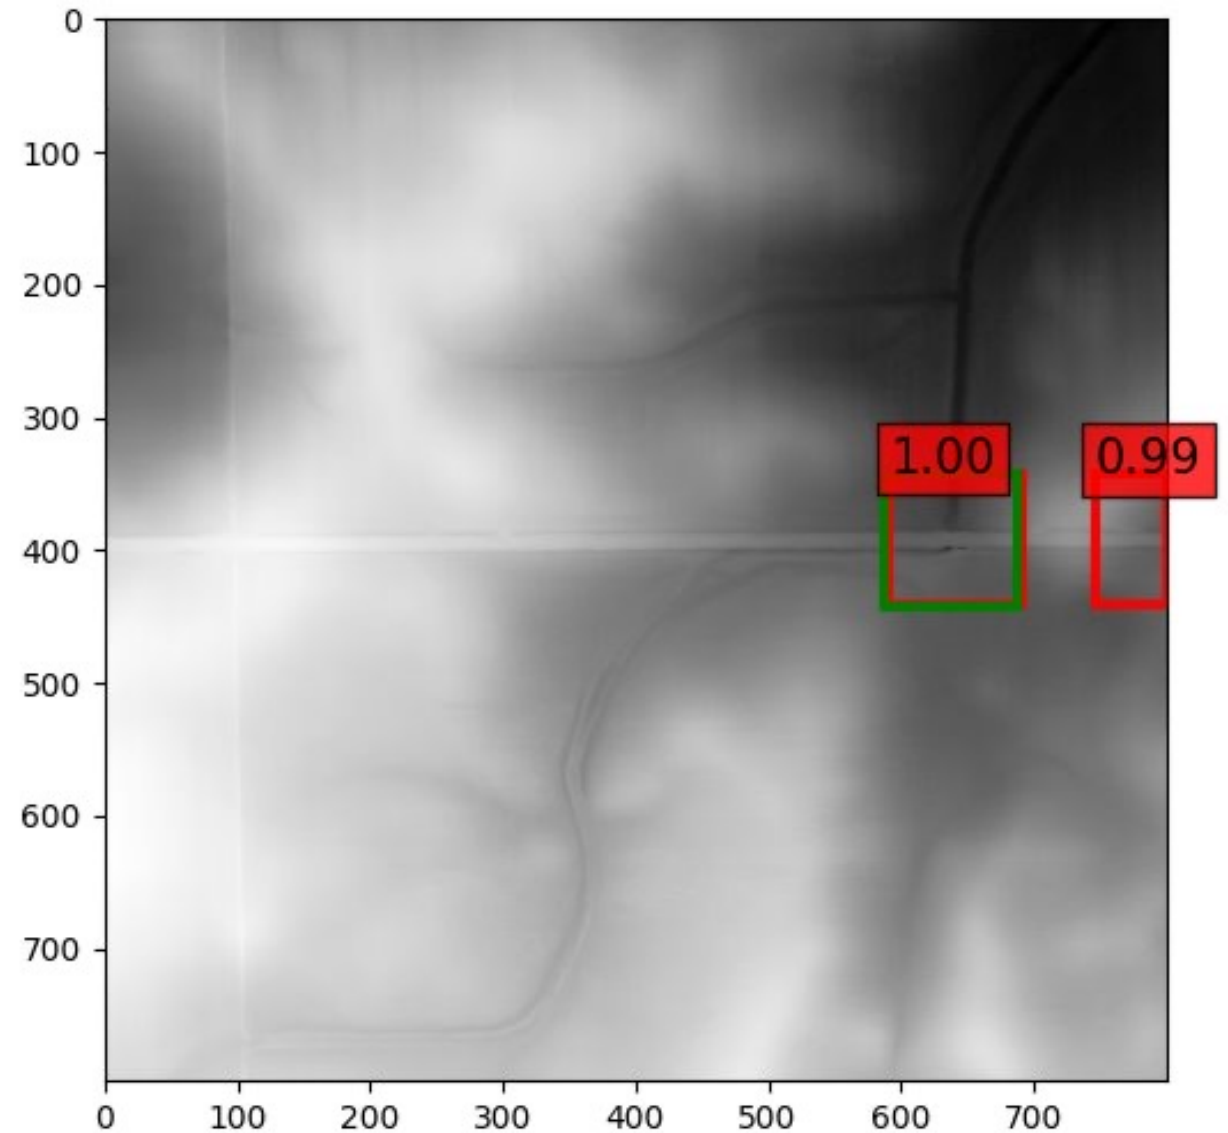
Case 3: Model finds an unambiguous culvert not labeled in the dataset



---

## MODEL 1 VISUALIZATIONS: PRETRAINED

Case 4: Human-in-the-loop training may be beneficial to distinguish ambiguous examples – the false positive here is an edge case




---

# MODEL 1 STATISTICS

Metric	Test Set Performance
Average Precision (AP) @[ IoU=0.50:0.95   area= all   maxDets=100 ]	0.645
Average Precision (AP) @[ IoU=0.50   area= all   maxDets=100 ]	0.788
Average Precision (AP) @[ IoU=0.75   area= all   maxDets=100 ]	0.747
Average Precision (AP) @[ IoU=0.50:0.95   area=medium   maxDets=100 ]	0.538
Average Precision (AP) @[ IoU=0.50:0.95   area= large   maxDets=100 ]	0.674
Average Recall (AR) @[ IoU=0.50:0.95   area= all   maxDets= 1 ]	0.322
Average Recall (AR) @[ IoU=0.50:0.95   area= all   maxDets= 10 ]	0.825
Average Recall (AR) @[ IoU=0.50:0.95   area= all   maxDets=100 ]	0.888
Average Recall (AR) @[ IoU=0.50:0.95   area=medium   maxDets=100 ]	0.828
Average Recall (AR) @[ IoU=0.50:0.95   area= large   maxDets=100 ]	0.904

# MODEL 2: FROM SCRATCH PERFORMANCE

---


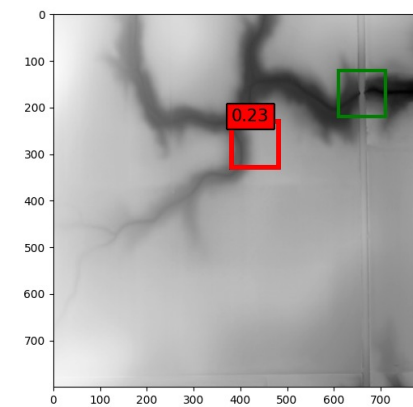
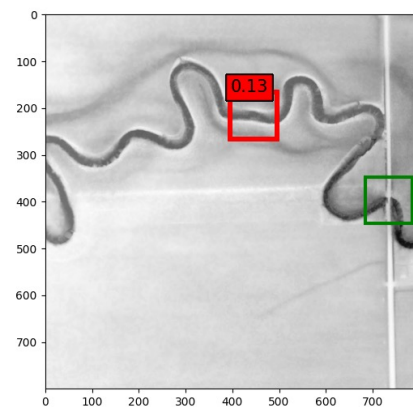










---

## MODEL 2 VISUALIZATIONS: FROM SCRATCH

- Training finds a bad local minimum – this training was repeated multiple times with similar results each time.
- Different hyperparameters and loading pre-trained CNN and encoder weights may yield better results for a more lightweight binary classifier.



---

# MODEL 2 STATISTICS

Metric	Test Set Performance
Average Precision (AP) @[ IoU=0.50:0.95   area= all   maxDets=100 ]	0.000
Average Precision (AP) @[ IoU=0.50   area= all   maxDets=100]	0.000
Average Precision (AP) @[ IoU=0.75   area= all   maxDets=100 ]	0.000
Average Precision (AP) @[ IoU=0.50:0.95   area=medium   maxDets=100 ]	0.000
Average Precision (AP) @[ IoU=0.50:0.95   area= large   maxDets=100 ]	0.000
Average Recall (AR) @[ IoU=0.50:0.95   area= all   maxDets= 1 ]	0.002
Average Recall (AR) @[ IoU=0.50:0.95   area= all   maxDets= 10 ]	0.002
Average Recall (AR) @[ IoU=0.50:0.95   area= all   maxDets=100 ]	0.002
Average Recall (AR) @[ IoU=0.50:0.95   area=medium   maxDets=100 ]	0.000
Average Recall (AR) @[ IoU=0.50:0.95   area= large   maxDets=100 ]	0.002

---

# SCALESIM

Systolic CNN AcceLErator Simulator

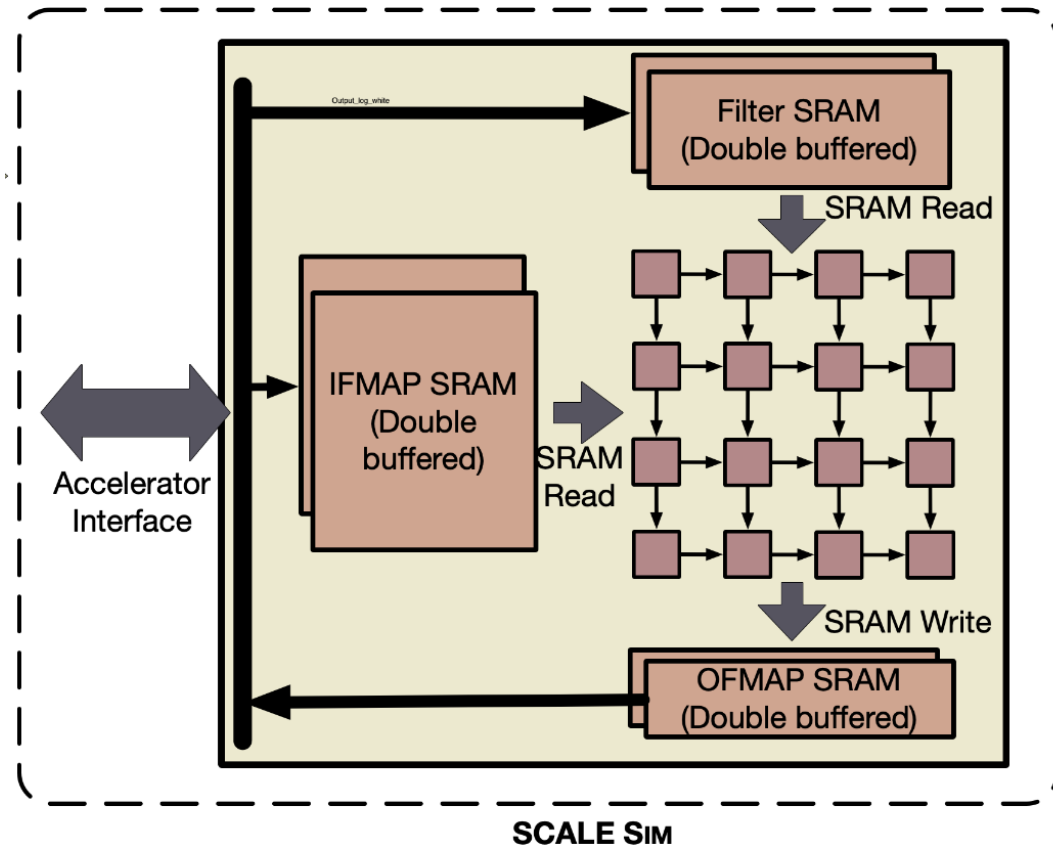
# USING SCALESIM FOR ARCHITECTURE SEARCH

## Configurable Parameters:

- Systolic Array Height and Width
- Memory size of Filter, IFMAP and OFMAP
- Bandwidth
- Number of Memory Banks
- Dataflow type (ws, os, is)

## Performance Related Output:

- Latency
- Mapping Efficiency
- Compute Utilization
- Per Layer



---

# DESIGN SPACE EXPLORATION

8 configurable parameters:

- 2 Array dimensions [260, 400]  $\rightarrow 140^2$  choices
- 3 SRAM sizes [1024, 35\*1024]  $\rightarrow 34816^3$  choices
- 3 types of Dataflow [ws, os, is]  $\rightarrow 3$  choices
- Number of Memory Banks [1, 4]  $\rightarrow 4$  choices
- Bandwidth [1, 64]  $\rightarrow 64$  choices
- Design Space size  $\approx 7e20$  points

ScaleSim experiment time:

- 800 x 800 input image  $\rightarrow 35$  minutes per experiment  $\rightarrow 35 * 7e20 = 174$  days for a greedy approach

---

# DESIGN SPACE EXPLORATION PRUNING

Solution:

- Limit range of available configuration values
- Use smaller input image size (260 x 260)
- Use random search or genetic algorithm for optimal design point

PyGAD:

- open-source Python Library for genetic algorithm optimization
- Custom fitness function, using one optimization objective, the total latency

---

# PYGAD GENETIC ALGORITHM

Genetic Algorithm properties:

- 50 generations
- 4 parent chromosomes
- 10 solutions per population
- 8 genes per chromosome

Chromosome Structure:

Array Height	Array Width	IFMAP size	Filter mem size	OFMAP size	Dataflow	Bandwidth	Memory Banks
-----------------	----------------	---------------	--------------------	---------------	----------	-----------	-----------------

---

# GENETIC ALGORITHM

## Chromosome Values

- Array Height and Width [260, 320]
- IFMAP, Filter SRAM [1 KB, 10 KB]
- OFMAP SRAM [512 bytes, 5 KB]
- Dataflow: ws, os, is
- Bandwidth bytes [1, 9]
- Memory Banks [1, 3]

---

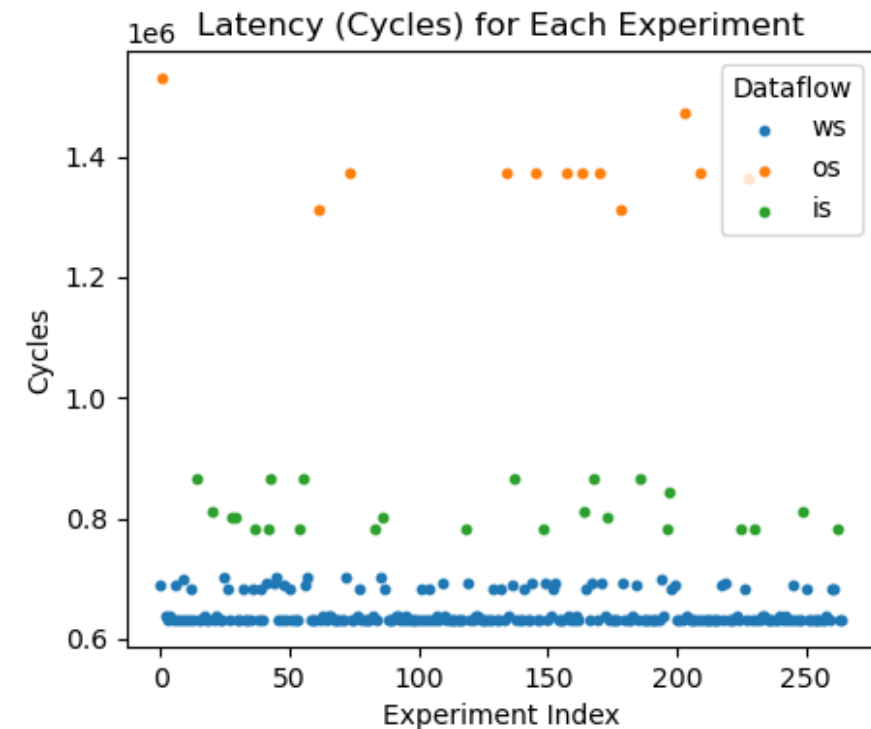
Converged after 2 days by executing 265 experiments

---


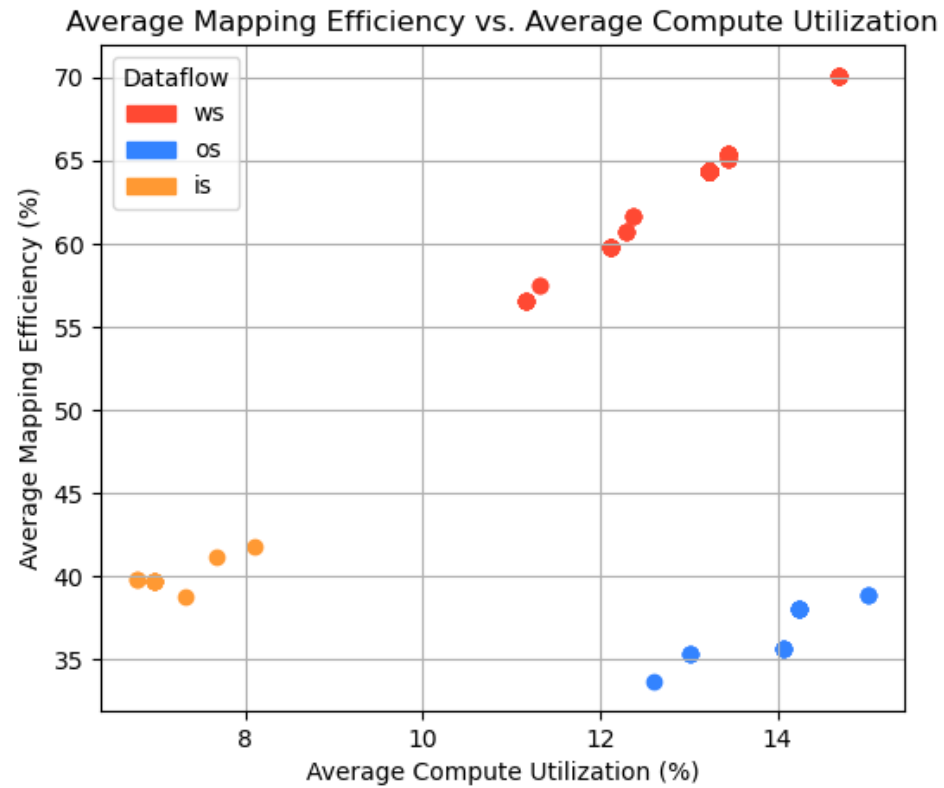


# OPTIMAL SOLUTION FOR 260 X 260 INPUT

Parameter	Optimal Value
Systolic Array dimensions	300 x 260
IFMAP SRAM size	1 KB
Filter SRAM size	2 KB
OFMAP SRAM size	3584 bytes
Dataflow	Weight Stationary
Bandwidth	9 bytes
Memory Banks	2
Total Latency	631451 cycles



# OPTIMAL TPU ARCHITECTURE



---

# ISSUES AND FUTURE WORK

1. Overtraining may occur prematurely as no random transforms are applied in training
2. The model is worse at detecting boxes in the edge of the image
3. The performant model processed 800x800 image chips, while the most efficient TPU architecture is found using a model with a 260x260 input size

Possible solution to all 3: implement random cropping as a training transform, cropping to 260x260 inputs, preventing overtraining and constraining model search parameters.

- This would most likely benefit from using a smaller bounding box size, e.g. 50x50
- Then, during inference, use a 260x260 sliding window with overlap to predict drainage crossing locations
- This is all possible in the current codebase with a fixed transforms script and retraining